

Combined 5-arm comparison -- 6 semantic-risk + 3 authority-risk scenarios
(controlled proof-of-mechanism, not a statistical benchmark -- see RESULTS.md)

